



# Lecture 6

## Introduction of Bayesian analysis using JAGS

WILD6900 (Spring 2020)

# Readings

| Kéry & Schaub 38-40; 58-64

# From custom MCMC to BUGS

# The BUGS language

## Bayesian Analysis Using Gibbs Sampling

Language/program invented in the 1990's by epidemiologists in Cambridge

Later became WinBUGS

- First customizable MCMC software
- Revolutionized the use of Bayesian methods in applied statistics

# The BUGS language

Since the development of WinBUGS, other Bayesian software programs have been developed:

- OpenBugs
- JAGS
- Nimble
- Stan

For the remainder of the course, we will fit models using JAGS

- Just **A**nother **G**ibbs **S**ampler
- Uses BUGS language (easy to learn, lots of online documentation)
- Often out-performs WinBUGS
- Available for all operating systems

# The BUGS language

Last week, we learned how to:

- Determine the *full conditional distributions* for a linear regression model
- Write a custom MCMC sampler to produce samples from the joint posterior distribution

Given a statistical model and user-specified prior distributions, JAGS does these steps for you!

- Possible to fit arbitrarily complex models <sup>1</sup>
- "Frees the modeler in you"

# Running JAGS from R

JAGS is a stand alone software program

- Can be run from GUI or command line

JAGS can also be run from R using the **jagsUI** package (among others)

- Write model in R script and save as **.jags** file
- Provide **jagsUI** with data, initial values, and MCMC settings
- model run in JAGS
- model output brought back in to R for diagnostics/analysis/visualization

# The BUGS language

Very similar to **R** (but more limited)

- Limited ability to vectorize operations

If you can write your model on the whiteboard,  
you can write it in JAGS

- Stochastic relationships represented by  $\sim$
- Deterministic relationships represented by  $\leftarrow$



# The BUGS language

## Linear regression model

$$y_i = \alpha + \beta \times x_i + \epsilon_i$$

$$\epsilon_i \sim \textit{Normal}(0, \tau)$$

# The BUGS language

## Linear regression model

$$y_i = \underbrace{\alpha + \beta \times x_i}_{\textit{Deterministic}} + \underbrace{\epsilon_i}_{\textit{Stochastic}}$$

# The BUGS language

## Linear regression model

$$\underbrace{\mu_i = \alpha + \beta \times x_i}_{\textit{Deterministic}}$$

$$\underbrace{y_i \sim \textit{Normal}(\mu_i, \tau)}_{\textit{Stochastic}}$$

Remember that these equations define the *likelihood* of our data given values of  $\alpha$ ,  $\beta$ , and  $\tau$

# The BUGS language

## Linear regression model

To specify a fully Bayesian model, we also need to define the priors:

$$[\alpha] \sim \text{Normal}(\alpha|0, 0.001)$$

$$[\beta] \sim \text{Normal}(\alpha|0, 0.001)$$

$$[\tau] \sim \text{Gamma}(\tau|0.01, 0.01)$$

# The BUGS language

## Linear regression model

```
model{  
  ## Priors  
  alpha ~ dnorm(0, 0.001)  
  beta ~ dnorm(0, 0.001)  
  tau ~ dgamma(.001,.001) # Precision  
  sigma <- 1/sqrt(tau)      # Calculate sd from precision  
  
  ## Likelihood  
  for(i in 1:N){  
    mu[i] <- alpha + beta * x[i]  
    y[i] ~ dnorm(mu[i], tau)  
  }  
} #end of model
```

# Writing model files

```
sink(file="jags/linear_regression.jags")
cat("
  model{
    ## Priors
    alpha ~ dnorm(0, 0.001)
    beta ~ dnorm(0, 0.001)
    tau ~ dgamma(.001,.001) # Precision
    sigma <- 1/sqrt(tau)      # Calculate sd from precision
    ## Likelihood
    for(i in 1:N){
      mu[i] <- alpha + beta * x[i]
      y[i] ~ dnorm(mu[i], tau)
    }
  } #end of model
  ", fill=TRUE)
sink()
```

# Preparing the data

```
## Read simulated data frame
dat <- readRDS("data/sim_seed_counts.rds")

## Store data for JAGS as list
jags_data <- list(y = dat$y, x = dat$visits.c, N = nrow(dat))

## Create function that returns random initial values
jags_inits <- function(){list(alpha = runif(1, 200, 300),
                               beta = runif(1, 25, 75),
                               tau = runif(1))}

## Tell JAGS which parameters we want to monitor
params <- c("alpha", "beta", "tau", "sigma")
```

# Run the model

*## Run the model*

```
jags_fit <- jagsUI::jags(data = jags_data, inits = jags_inits,  
                        parameters.to.save = params,  
                        model.file = "jags/linear_regression.jags",  
                        n.chains = 3, n.iter = 10000,  
                        n.burnin = 2500, n.thin = 1)
```



# Diagnostics

```
print(jags_fit)
```

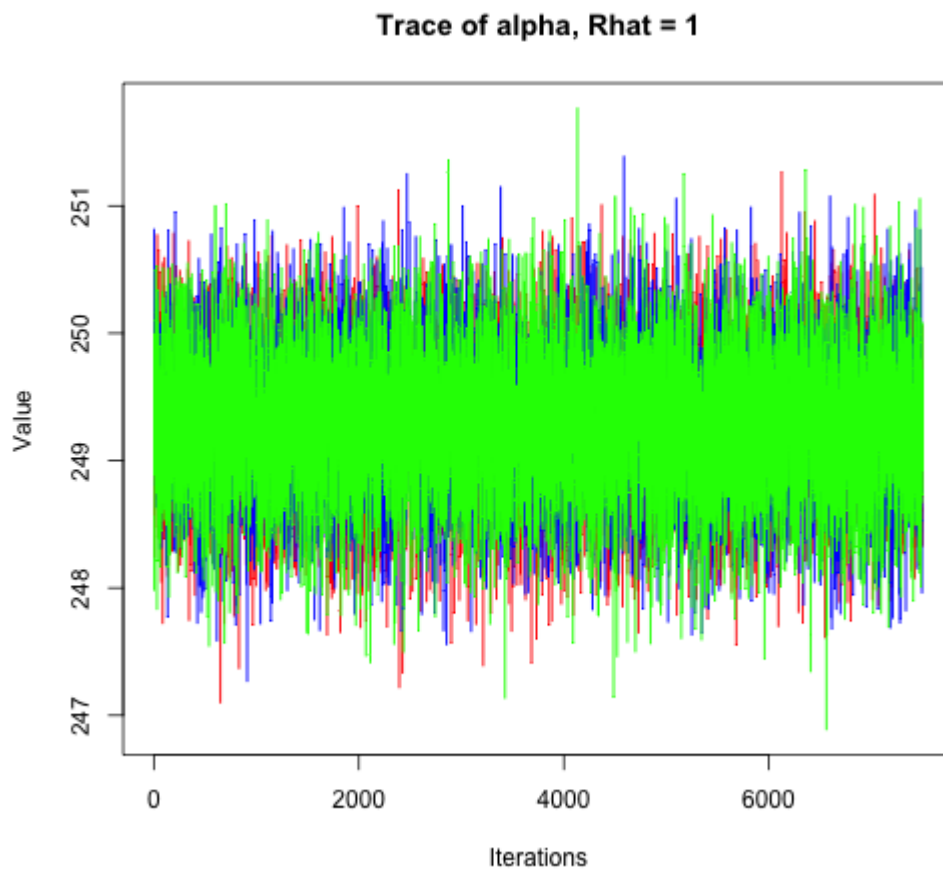
```
## JAGS output for model '/Library/Frameworks/R.framework/Versions/3.6/Resources/Model
## Estimates based on 3 chains of 10000 iterations,
## adaptation = 100 iterations (sufficient),
## burn-in = 2500 iterations and thin rate = 1,
## yielding 22500 total samples from the joint posterior.
## MCMC ran for 0.026 minutes at time 2020-01-04 20:36:37.
```

```
##
##           mean      sd      2.5%      50%      97.5% overlap0 f  Rhat n.eff
## alpha      249.300 0.551  248.198  249.302  250.367    FALSE 1 1.000 22500
## beta       50.104 0.558   48.999   50.104   51.184    FALSE 1 1.000 22500
## tau        0.019 0.002   0.015   0.019   0.023    FALSE 1 1.000 22500
## sigma      7.321 0.395   6.594   7.305   8.143    FALSE 1 1.000 22500
## deviance 1192.708 2.505 1189.884 1192.051 1199.172    FALSE 1 1.001  8610
##
```

```
## Successful convergence based on Rhat values (all < 1.1).
## Rhat is the potential scale reduction factor (at convergence, Rhat=1).
## For each parameter, n.eff is a crude measure of effective sample size.
```

# Diagnostics

```
jagsUI::traceplot(jags_fit)
```



# Structure of the JAGS output

```
class(jags_fit)
```

```
## [1] "jagsUI"
```

```
names(jags_fit)
```

```
## [1] "sims.list" "mean" "sd" "q2.5" "q25"  
## [6] "q50" "q75" "q97.5" "overlap0" "f"  
## [11] "Rhat" "n.eff" "pD" "DIC" "summary"  
## [16] "samples" "modfile" "model" "parameters" "mcmc.info"  
## [21] "run.date" "parallel" "bugs.format" "calc.DIC"
```

# Structure of the JAGS output

```
str(jags_fit$sims.list)
```

```
## List of 5
## $ alpha    : num [1:22500] 249 249 249 250 249 ...
## $ beta     : num [1:22500] 50.4 49.3 50.9 51.2 50.7 ...
## $ tau      : num [1:22500] 0.0218 0.0216 0.0206 0.0193 0.0208 ...
## $ sigma    : num [1:22500] 6.77 6.8 6.97 7.2 6.94 ...
## $ deviance: num [1:22500] 1192 1194 1192 1195 1191 ...
```

```
head(jags_fit$sims.list$alpha)
```

```
## [1] 249.4 249.5 249.2 250.1 249.3 248.2
```

# Structure of the JAGS output

```
jags_fit$mean$alpha
```

```
## [1] 249.3
```

```
jags_fit$f$alpha
```

```
## [1] 1
```

# Structure of the JAGS output

```
jags_fit$summary
```

```
##              mean      sd      2.5%      25%      50%      75%
## alpha      2.493e+02 0.551199 2.482e+02 2.489e+02 2.493e+02 2.497e+02
## beta       5.010e+01 0.557579 4.900e+01 4.973e+01 5.010e+01 5.048e+01
## tau        1.882e-02 0.002019 1.508e-02 1.744e-02 1.874e-02 2.013e-02
## sigma      7.321e+00 0.394910 6.594e+00 7.048e+00 7.305e+00 7.573e+00
## deviance   1.193e+03 2.505359 1.190e+03 1.191e+03 1.192e+03 1.194e+03
##              97.5%  Rhat n.eff overlap0 f
## alpha      250.367 1.000 22500          0 1
## beta       51.184 1.000 22500          0 1
## tau        0.023 1.000 22500          0 1
## sigma      8.143 1.000 22500          0 1
## deviance 1199.172 1.001  8610          0 1
```

# Structure of the JAGS output

```
str(jags_fit$samples)
```

```
## List of 3
## $ : 'mcmc' num [1:7500, 1:5] 249 249 249 250 249 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:5] "alpha" "beta" "tau" "sigma" ...
## ..- attr(*, "mcpair")= num [1:3] 2501 10000 1
## $ : 'mcmc' num [1:7500, 1:5] 249 251 249 249 249 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:5] "alpha" "beta" "tau" "sigma" ...
## ..- attr(*, "mcpair")= num [1:3] 2501 10000 1
## $ : 'mcmc' num [1:7500, 1:5] 249 248 250 250 249 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:5] "alpha" "beta" "tau" "sigma" ...
## ..- attr(*, "mcpair")= num [1:3] 2501 10000 1
## - attr(*, "class")= chr "mcmc.list"
```

# Structure of the JAGS output

```
str(jags_fit$mcmc.info)
```

```
## List of 9
## $ n.chains      : num 3
## $ n.adapt       : num 100
## $ sufficient.adapt: logi TRUE
## $ n.iter        : num 10000
## $ n.burnin      : num 2500
## $ n.thin        : num 1
## $ n.samples     : num 22500
## $ end.values    :List of 3
## ..$ : Named num [1:5] 2.49e+02 5.04e+01 1.19e+03 7.85 1.62e-02
## .. ..- attr(*, "names")= chr [1:5] "alpha" "beta" "deviance" "sigma" ...
## ..$ : Named num [1:5] 2.50e+02 5.07e+01 1.19e+03 7.62 1.72e-02
## .. ..- attr(*, "names")= chr [1:5] "alpha" "beta" "deviance" "sigma" ...
## ..$ : Named num [1:5] 2.49e+02 5.10e+01 1.19e+03 7.31 1.87e-02
## .. ..- attr(*, "names")= chr [1:5] "alpha" "beta" "deviance" "sigma" ...
## $ elapsed.mins   : num 0.026
```



# Saving model output

```
saveRDS(jags_fit, "output/regression_out.rds")
```